



Viller Hsiao <villerhsiao@gmail.com>

F9 Microkernel ktimer

Rights to Copy

copyright © 2014 Viller Hsiao



Attribution - ShareAlike 3.0

You are free

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Under the following conditions

- Ⓒ **BY:** – **Attribution.** You must give the original author credit.
- Ⓒ **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.

For any reuse or distribution, you must make clear to others the license terms of this work.

- Any of these conditions can be waived if you get permission from the copyright holder.

Your fair use and other rights are in no way affected by the above.

License text: <http://creativecommons.org/licenses/by-sa/3.0/legalcode>



Corrections, suggestions,
contributions and
translations are welcome!

F9 Microkernel [1]



- New open source implementation built from scratch, which deploys modern kernel techniques, derived from L4 microkernel designs, to deep embedded devices.
 - BSD License
- Characteristics of F9 microkernel
 - Efficiency: performance + power consumption
 - Tickless
 - Security: memory protection + isolated execution
 - Flexible development environment
- Repository
 - <https://github.com/f9micro>

Agenda



- Concept about operating system kernel time subsystem
- F9 Microkernel ktimer introduction



Concept about Operating System Kernel Time Subsystem

OS Kernel Time Subsystem



- Current Time
 - System time, UTC/GMP time, ... etc
- Alarm, the time event
- Process timeslice
- Timeout
- Time delay

System Time and Tick



- From Wikipedia [2],


“System time represents a computer system's notion of the passing of time.”

“System time is measured by a system clock, which is typically implemented as a simple count of the number of ticks that have transpired since some arbitrary starting date, called the epoch.”

How Is Tick Implemented?




- Hardware timer device
 - Assert interrupt after a programmable interval
 - Handling tick stuff in Timeout Interrupt Service Routine (ISR)



How Is Tick Implemented? (Cont.)



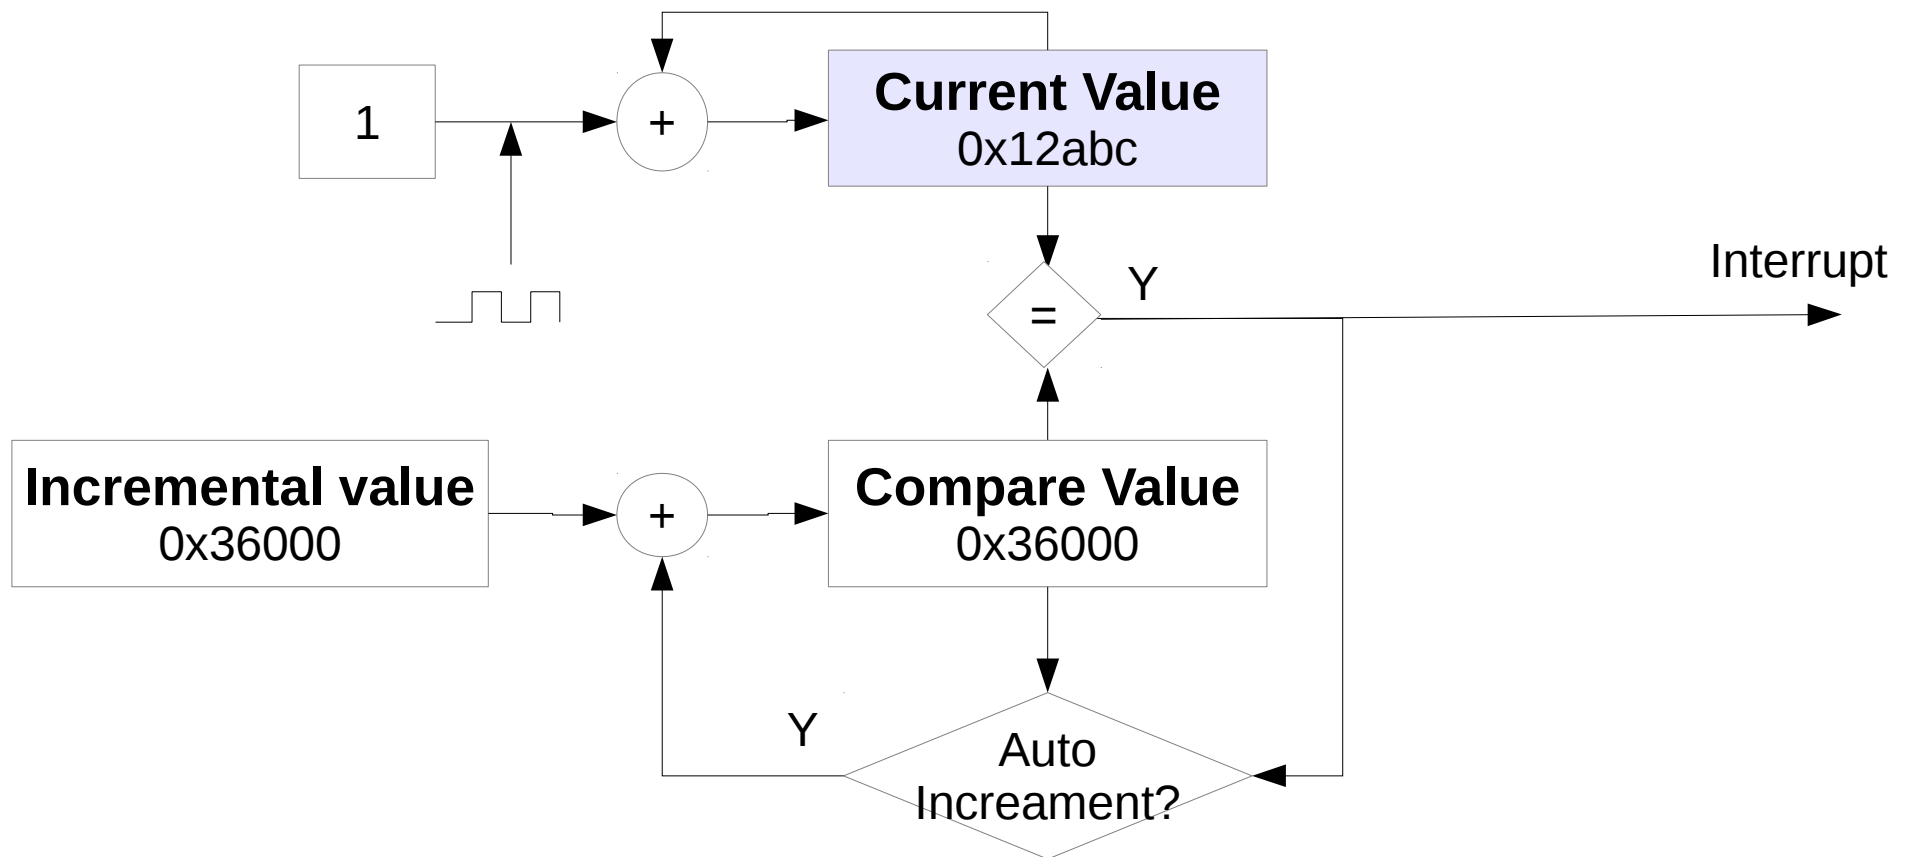
- Type 1: Incremental timer
 - Example: Cortex-A9 MP Global Timer





How Is Tick Implemented? (Cont.)

- Incremental timer functionality



How Is Tick Implemented? (Cont.)



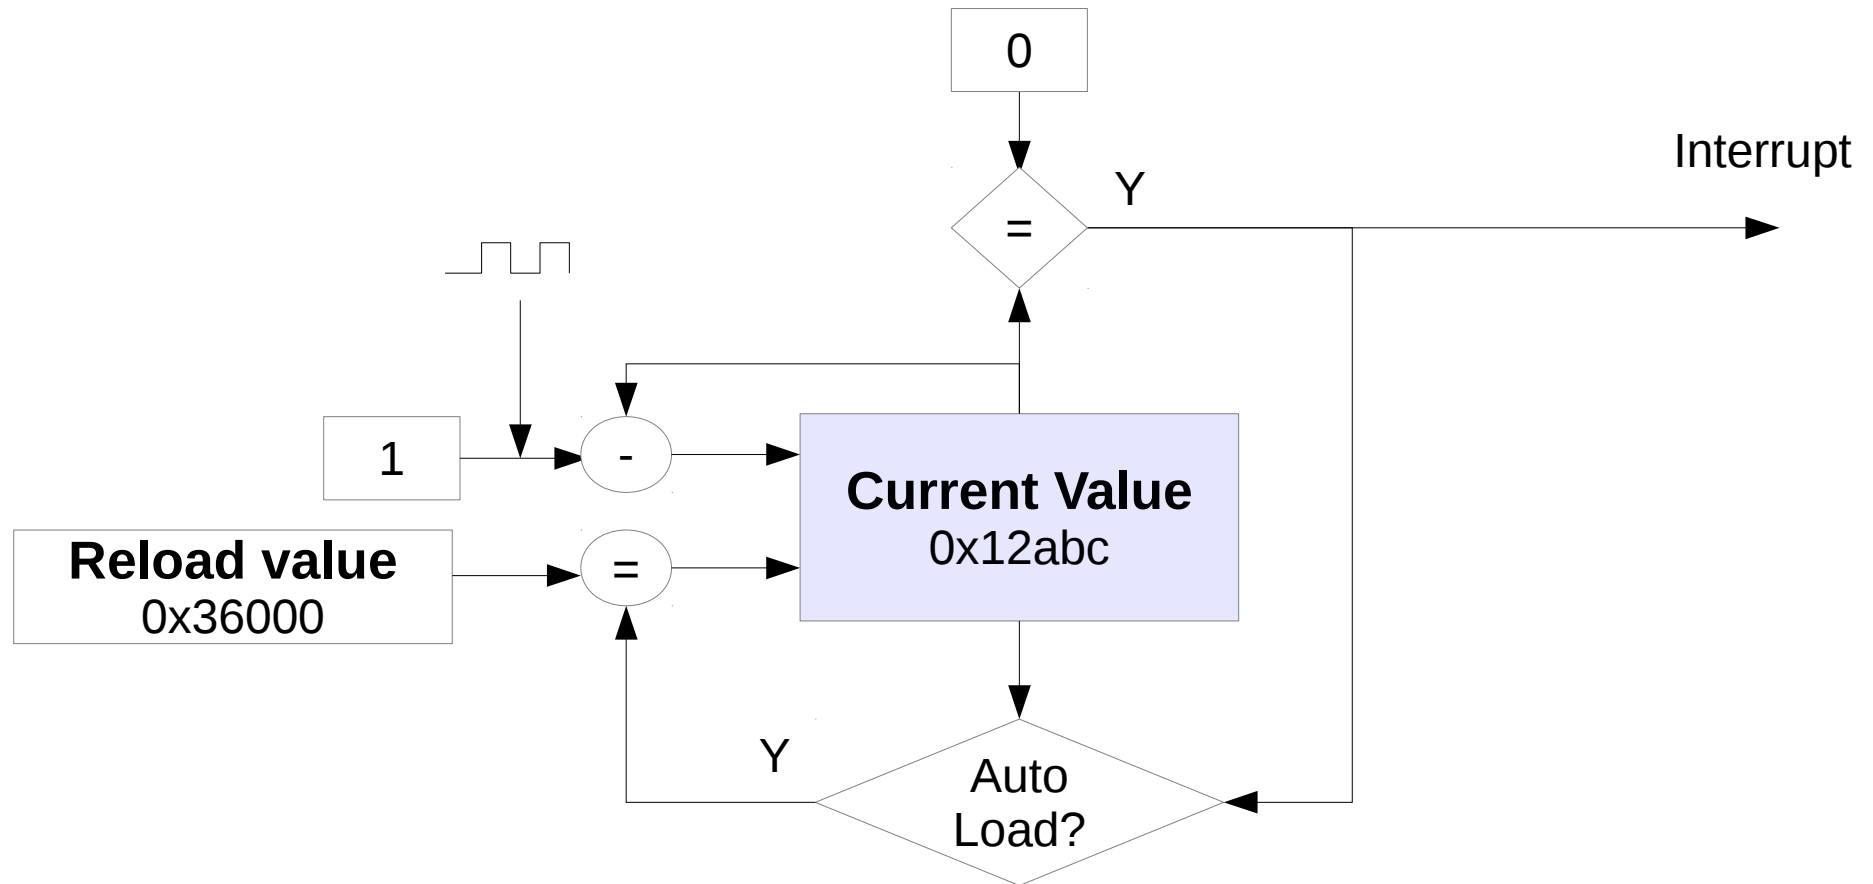
- Type 2: Counting down timer
 - Example: Cortex-M4 SysTick





How Is Tick Implemented? (Cont.)

- Counting down timer functionality

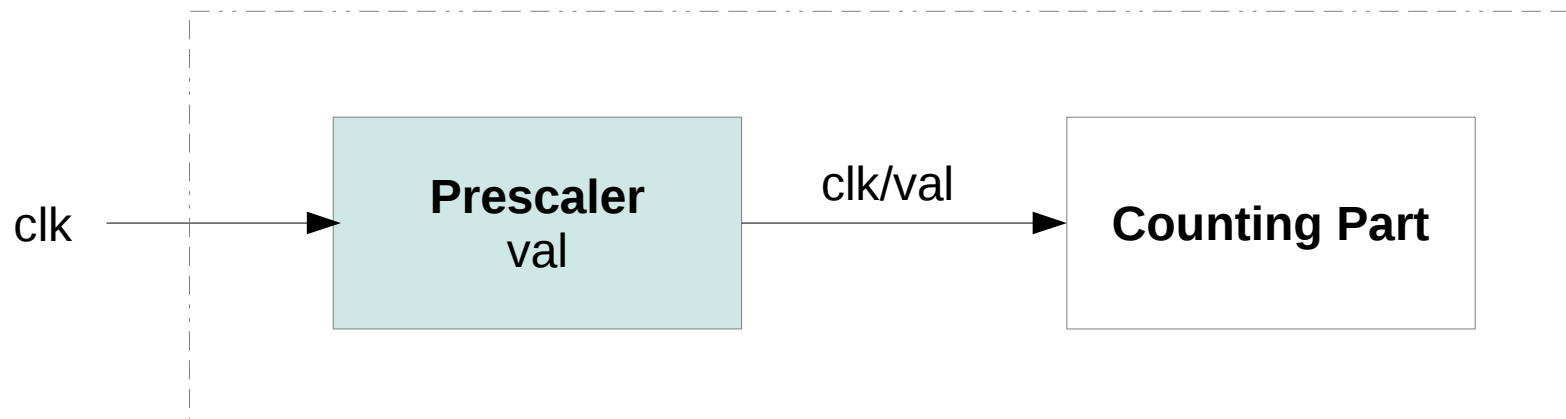


How Is Tick Implemented? (Cont.)



- Prescaler
 - Tweak HW tick period

Timer Module

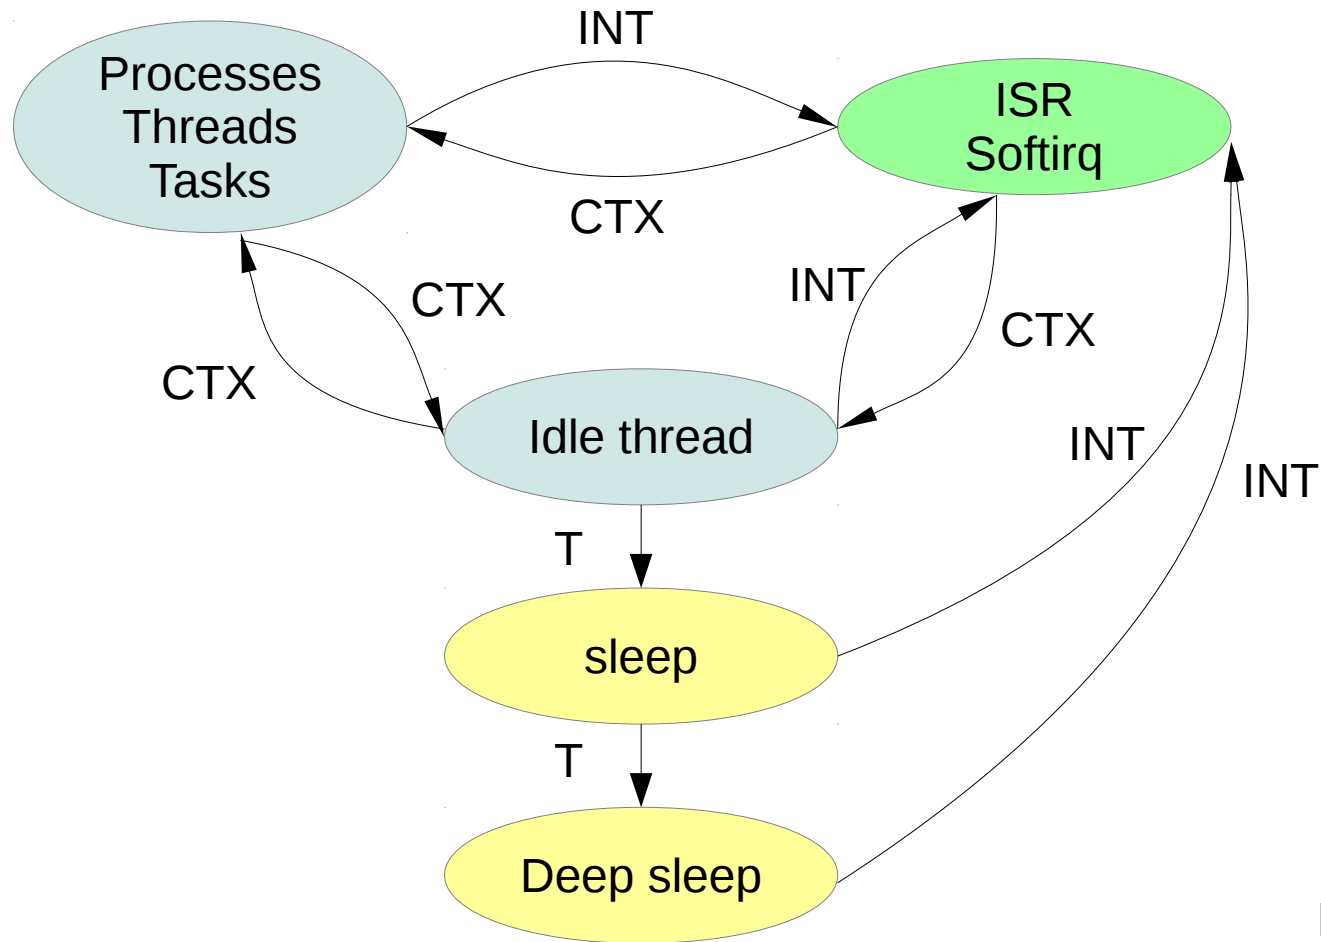


How Is Tick Implemented? (cont.)



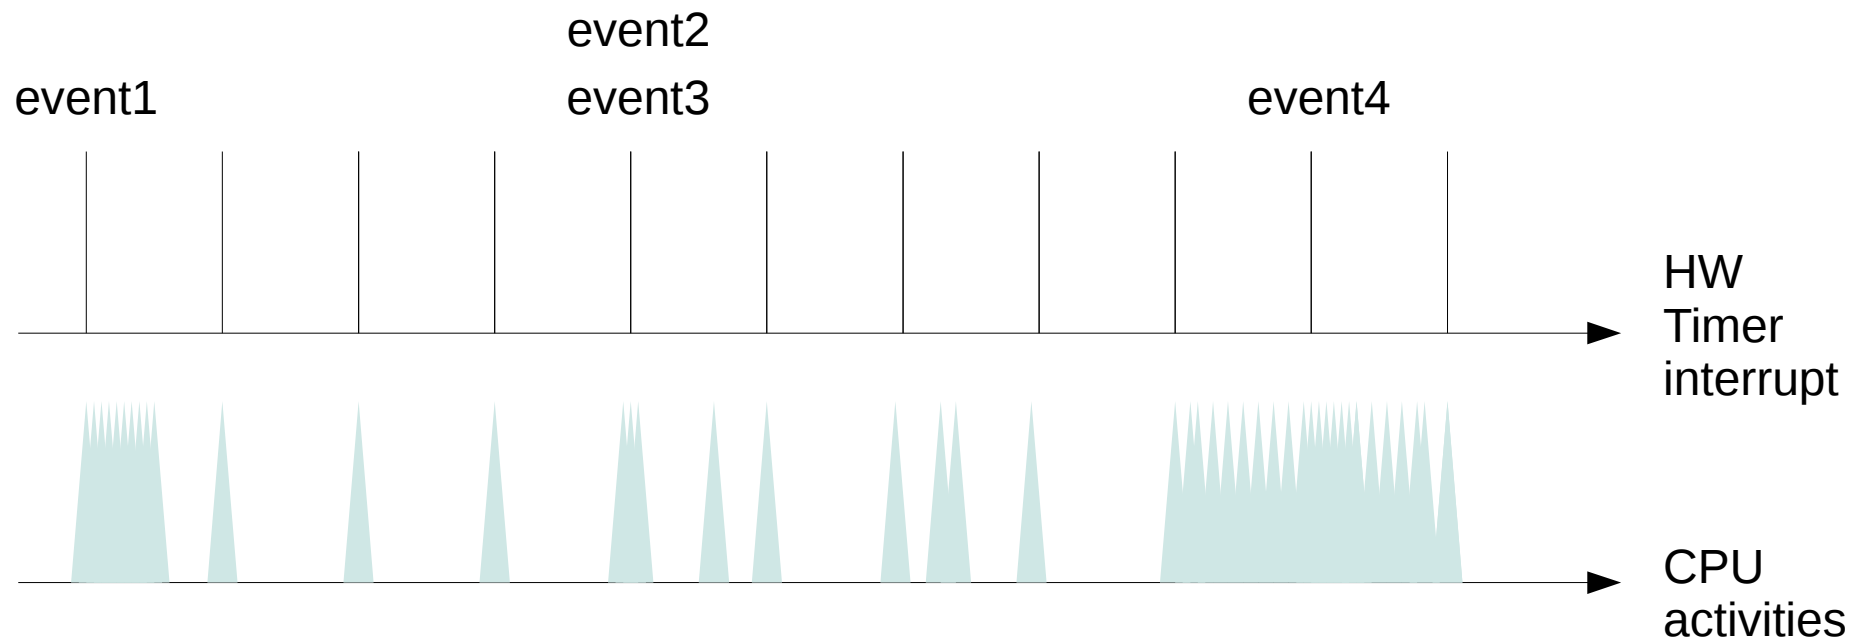
- Timeout ISR
 - Increase system ticks
 - Execute handler of timeout event
 - Re-schedule if required

Simple Graph of CPU Operating States



INT : interrupt
CTX : context switch
T : after a while

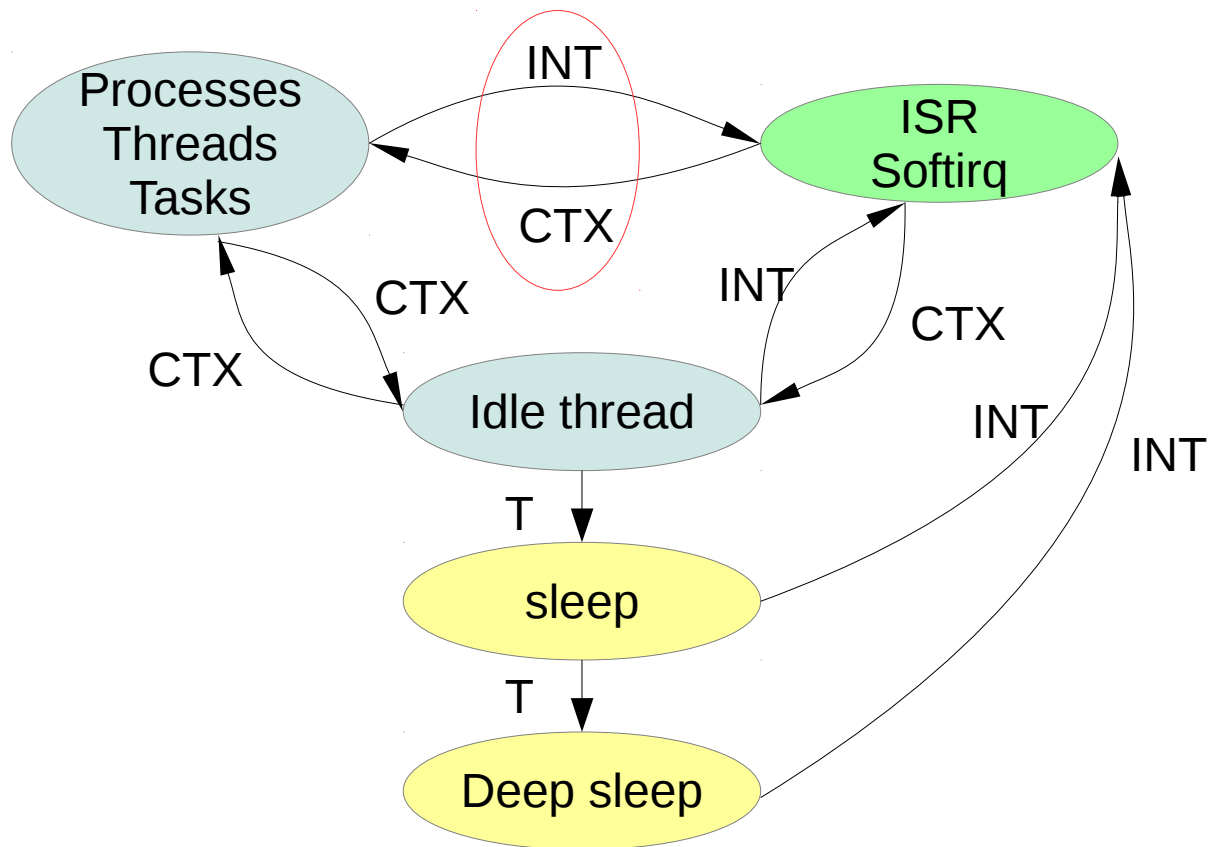
Time Diagram of Legacy Ticks



Where Is the Problem?



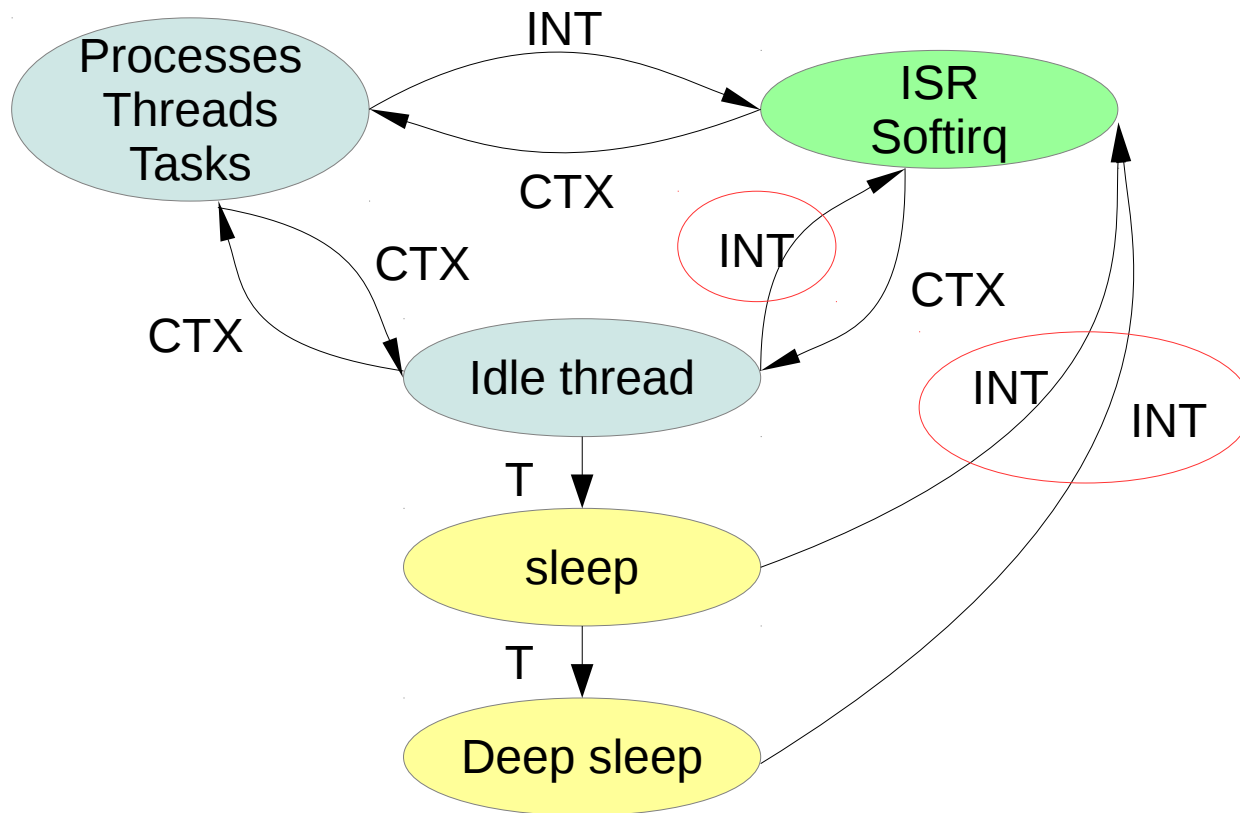
- Context switch overhead



Where Is the Problem?



- Unnecessary power consumed

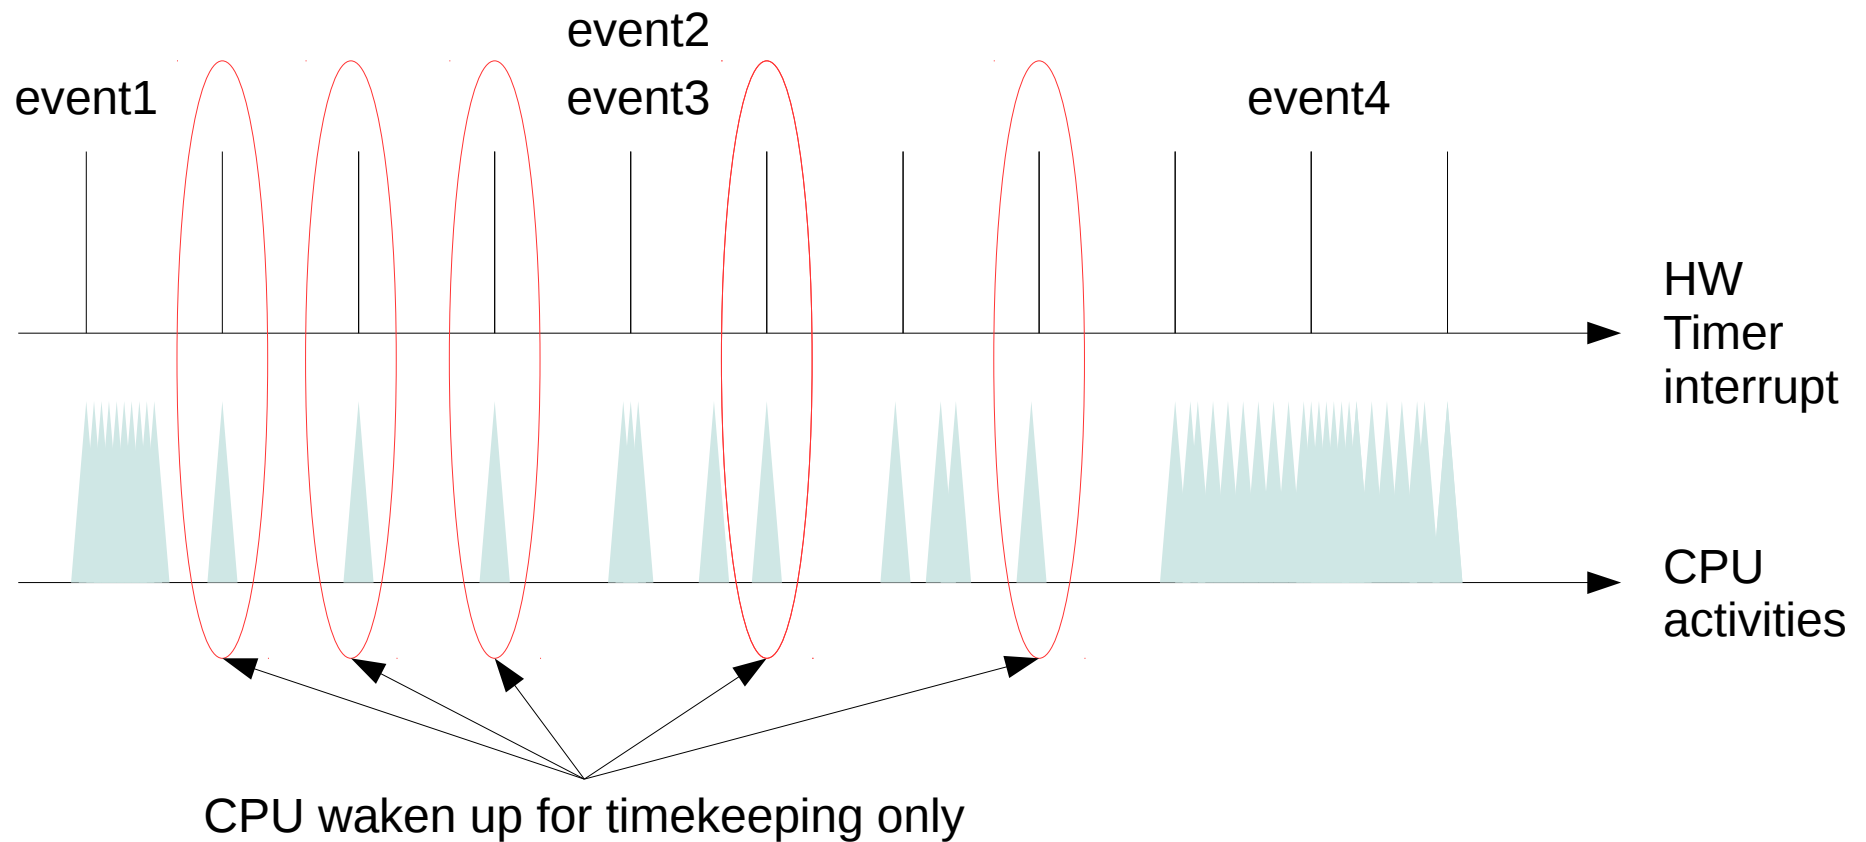


Where Is the Problem?




- Tick resolution
 - Performance, power v.s. timer precision

Time Diagram of Legacy Ticks



Solution: Tickless



Drawback of Tickless



- Tickless is not free [3],
“It increases the number of instructions executed on the path to and from the idle loop.”
“On many architectures, dyntick-idle mode also increases the number of expensive clock-reprogramming operations”
- Systems with aggressive real-time response constraints often run periodic tick

Example: Linux Time Subsystem



- Device level
 - clocksource
 - jiffies
 - clockevent
- High resolution kernel timer
 - hrtimers

Example: Linux Time Subsystem (Cont.)



- Tickless
 - (Nearly) Full tickless operation after 3.10
 - **dyntick-idle**
 - CONFIG_NO_HZ_IDLE
 - **adaptive-tick**
 - CONFIG_NO_HZ_FULL
 - **periodic**
 - CONFIG_HZ_PERIODIC
 - Consumes battery power 2-3 times as fast as **dyntick-idle** one. [3]



F9 Microkernel ktimer Introduction

Example Hardware Device



- STM32F4Discovery board
 - Cortex-M4 EVB
 - Use Cortex-M4 SysTick as the tick HW
 - 24-bit system timer
 - Counting down
 - Please refer to [4] for more detail

ktimer Features



- System time
- Time event
- Tickless

ktimer Configurations



- CONFIG_KTIMER_HEARTBEAT
 - HW cycles per ktimer tick
- CONFIG_KTIMER_MINTICKS
 - Minimum ktimer ticks unit for time event
- CONFIG_KTIMER_TICKLESS
 - Enable tickless

ktimer Event Utility



- **ktimer_event_create()**
 - Arguments
 - Timeout ticks
 - Timeout handler
 - Pointer to arguments for timeout handler
 - Return Non-zero to reschedule event again
 - Freed automatically if not be re-scheduled

ktimer Anatomy



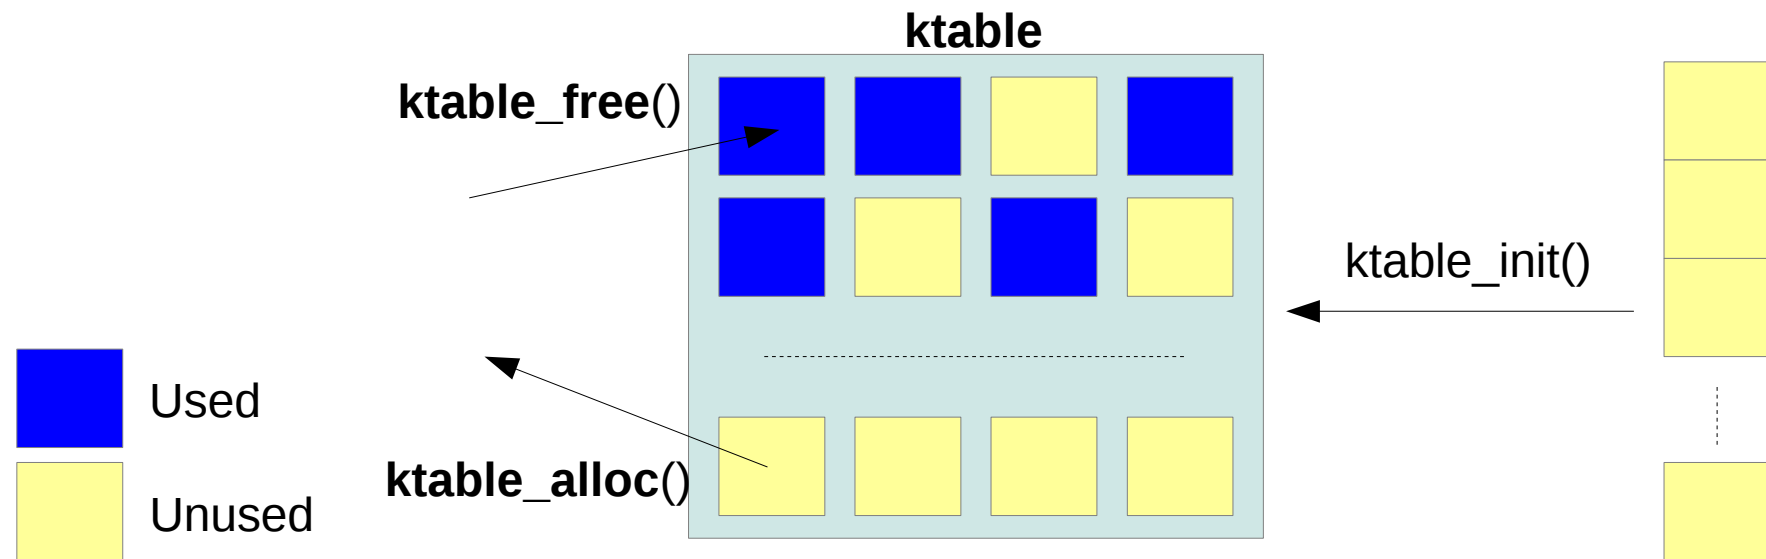
- System time
 - **ktime_now**
- ktimer event list
 - **event_queue**
 - Managed by **ktable**

ktimer Anotomy (Cont.)



- **htable**

- Basic kernel library in F9 Microkernel
- Fast objects pool management
 - To allocate/free objects of pre-defined size and numbers easiler



ktimer Anatomy (Cont.)



- Life cycle of time event
 - One-shot
 - Continuous
- How to control life cycle in ktimer?
 - The return value of handler is treated as next timeout delta

ktimer Anotomy (Cont.)



- HW Timer ISR
 - **__ktimer_handler()**
 - Increase tick
 - Arrange Softirq for handling timer events

ktimer Anatomy (Cont.)



- Two-stage interrupt handling
 - ISR
 - IRQ context
 - Softirq
 - Thread context
 - Real time preemptive characteristic.
 - Can be scheduled like any other threads in the system [5].
 - Handled in **kernel thread**

ktimer Anotomy (Cont.)



- HW Timer ISR
 - **__ktimer_handler()**
 - Increase tick
 - Arrange Softirq for handling timer events

ktimer Anotomy (Cont.)



- HW timer softirq handler
 - **ktimer_event_handler()**
 - Find timeout event and executing its handler
 - Re-insert event if required

ktimer Anatomy (Cont.)

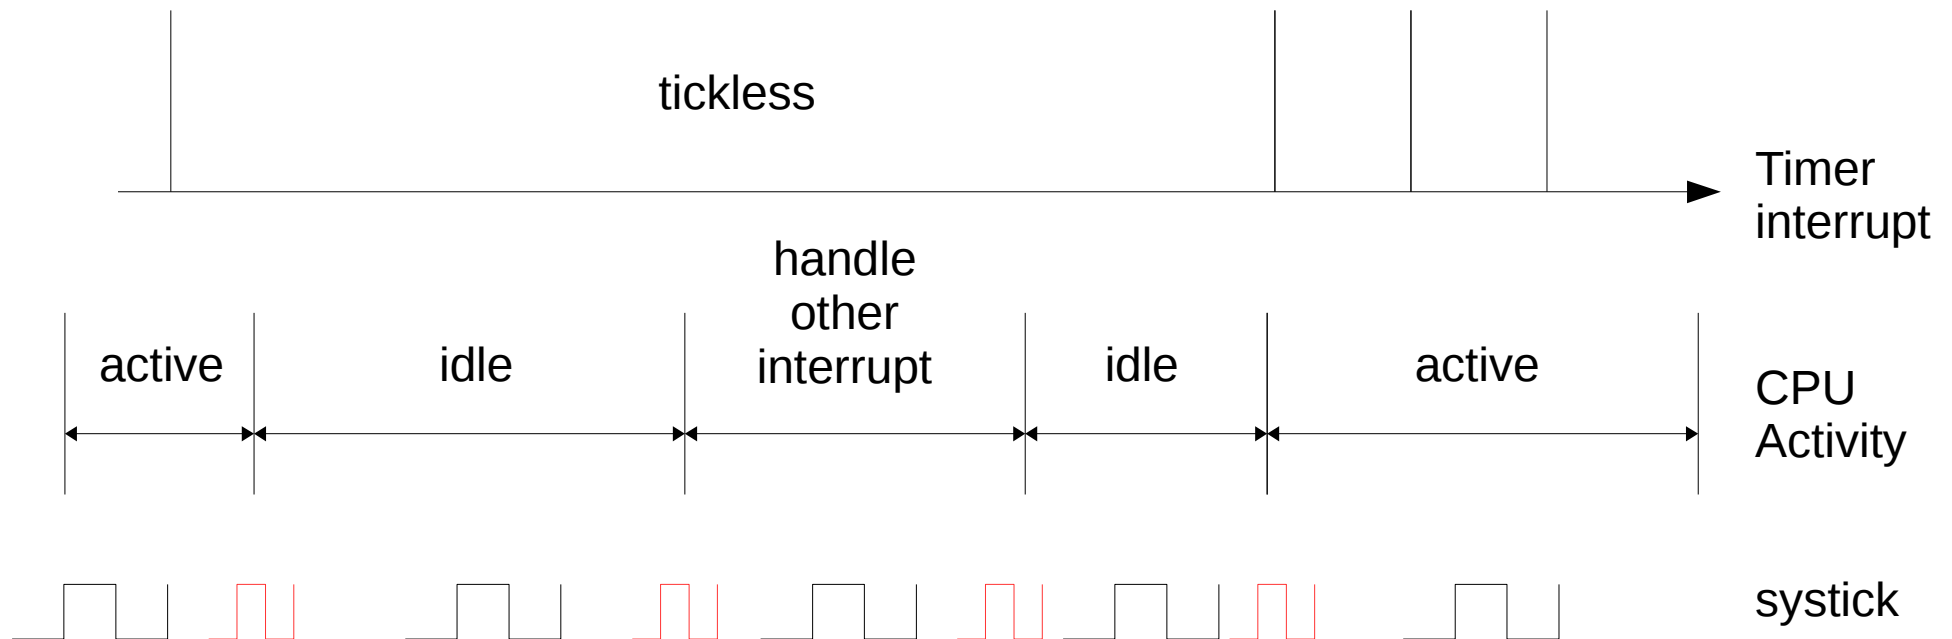


- Tickless
 - Enter tickless right before going to CPU idle state
 - Set interval of next timer interrupt as delta of next event
 - Or **KTIMER_MAXTICKS**
 - Adjust system time after waked up

ktimer Anatomy (Cont.)




- Tickless Compensation
 - SysTick frequency distortion when enter/exit standby mode



ktimer Anatomy (Cont.)



- Tickless Compensation



ktimer Anatomy (Cont.)




- Tickless Compensation
 - Pre-defined value
 - **CONFIG_KTIMER_TICKLESS_COMPENSATION**
 - Compensation value before entering idle
 - **CONFIG_KTIMER_TICKLESS_INT_COMPENSATION**
 - Compensation value when CPU is waked up by non-timer interrupt.
 - But how much to set these value to?

ktimer Anatomy (Cont.)



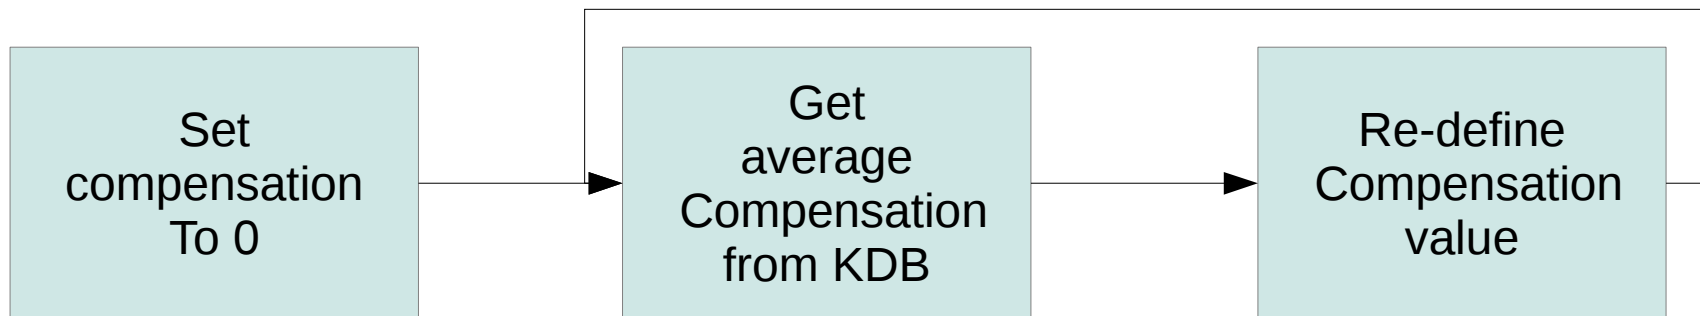
- Tickless Compensation
 - But how much to set these value to?
 - Get average compensation from general pupose timer
 - $(\text{total diff} / \text{count of entering tickless})$



ktimer Anotomy (Cont.)



- Tickless Compensation
 - How to calibrate the compensation value?
 - **CONFIG_KTIMER_TICKLESS_VERIFY**
 - **CONFIG_KDB**



ktimer Anotomy (Cont.)



- KDB
 - Kernel debugger

```
## KDB ##
commands:
K: print kernel tables
e: dump ktimer events
n: show timer (now)
s: show softirqs
t: dump threads
m: dump memory pools
a: dump address spaces
p: show sampling
v: show tickless scheduling stat
-----
```

```
## KDB ##
-----NOW-----
Now is 70993
Ktimer T=1297 D=0
-----
```

Reference:



- [1] Jim Huang (Dec 9, 2013),
“F9: A Secure and Efficient Microkernel Built for Deep Embedded Systems”
- [2] Wikipedia, “System time”
- [3] P. E. McKenney (May 14, 2013), “NO_HZ: Reducing Scheduling-Clock Ticks”
- [4] ST Microelectronics,
STM32F3xxx and STM32F4xxx Cortex-M4 programming manual
- [5] J. Altenberg, “Using the Realtime Preemption Patch on ARM CPUs”

THE END